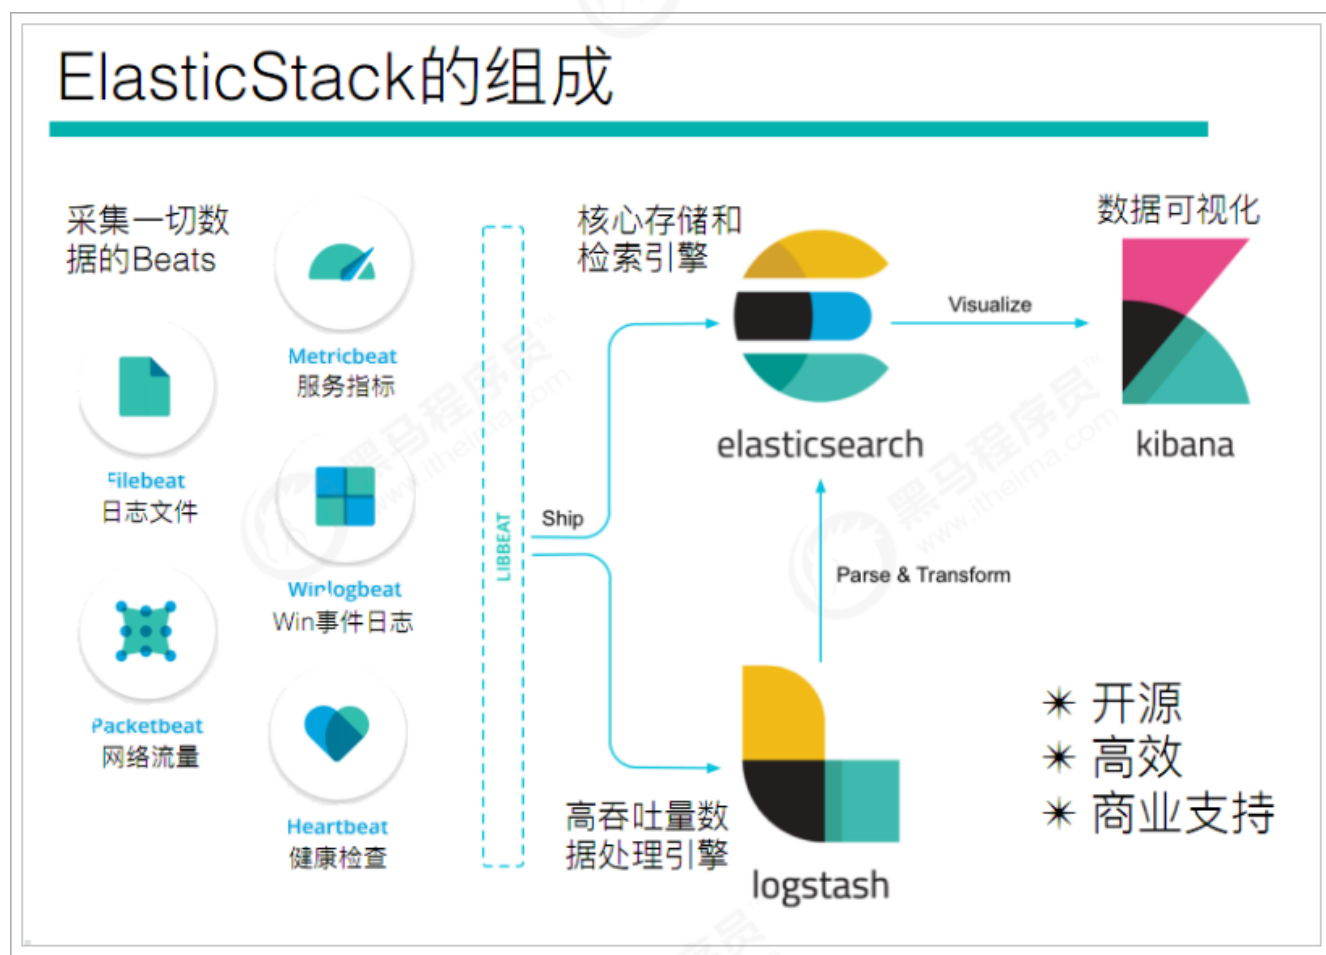


课程介绍

- 了解Beats
- Filebeat入门学习
- Metricbeat入门学习
- Kibana入门学习
- Logstash入门学习

1、Beats 简介



官网：<https://www.elastic.co/cn/products/beats>



轻量型数据采集器

Beats 平台集合了多种单一用途数据采集器。它们从成百上千或成千上万台机器和系统向 Logstash 或 Elasticsearch 发送数据。

Beats系列产品：

Beats 系列

全品类采集器，搞定所有数据类型。



Filebeat
日志文件



Metricbeat
指标



Packetbeat
网络数据




Winlogbeat
Windows 事件日志



Auditbeat
审计数据



Heartbeat
运行时间监控



Functionbeat
无需服务器的采集器

2、Filebeat



Filebeat

轻量型日志采集器

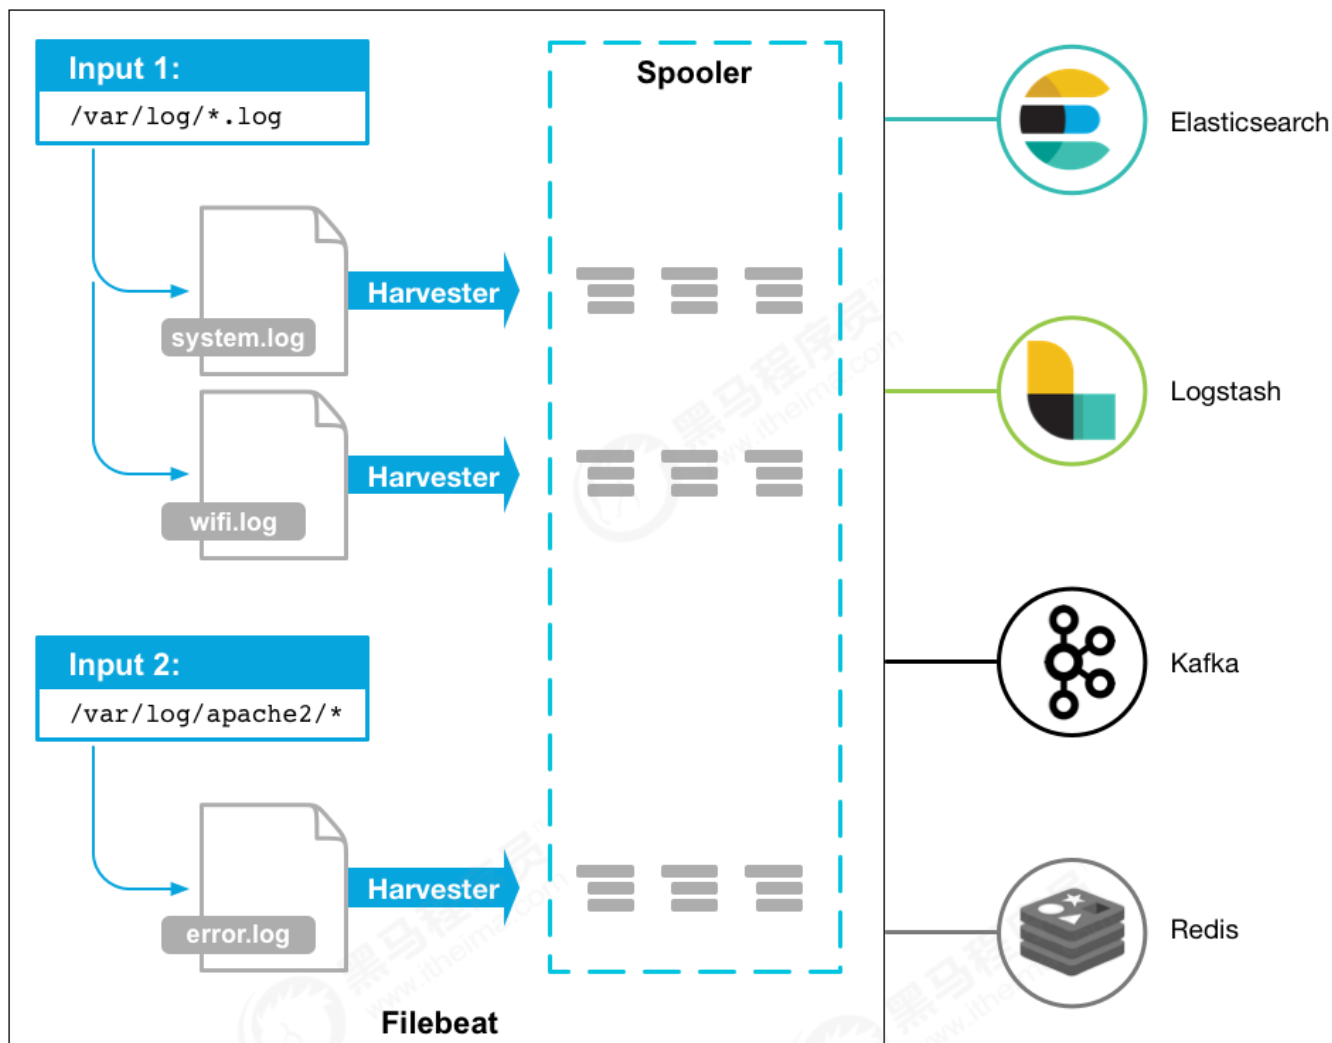
当您要面对成百上千、甚至成千上万的服务器、虚拟机和容器生成的日志时，请告别 SSH 吧。Filebeat 将为您提供一种轻量型方法，用于转发和汇总日志与文件，让简单的事情不再繁杂。

汇总、“tail -f” 和搜索

启动 Filebeat 后，打开 Logs UI，直接在 Kibana 中观看对您的文件进行 tail 操作的过程。通过搜索栏按照服务、应用程序、主机、数据中心或者其他条件进行筛选，以跟踪您的全部汇总日志中的异常行为。

2.1、架构

用于监控、收集服务器日志文件。



2.2、部署与运行

下载（或使用资料中提供的安装包，版本为：filebeat-6.5.4）：<https://www.elastic.co/downloads/beats>

```
1 mkdir /haoke/beats
2 tar -xvf filebeat-6.5.4-linux-x86_64.tar.gz
3 cd filebeat-6.5.4-linux-x86_64
4
5 #创建如下配置文件 haoke.yml
6 filebeat.inputs:
7   - type: stdin
8     enabled: true
9   setup.template.settings:
10     index.number_of_shards: 3
11   output.console:
12     pretty: true
13     enable: true
14
15 #启动filebeat
16 ./filebeat -e -c haoke.yml
17
18 #输入hello运行结果如下：
19 hello
```



```
20 {
21   "@timestamp": "2019-01-12T12:50:03.585z",
22   "@metadata": { #元数据信息
23     "beat": "filebeat",
24     "type": "doc",
25     "version": "6.5.4"
26   },
27   "source": "",
28   "offset": 0,
29   "message": "hello", #输入的内容
30   "prospector": { #标准输入勘探器
31     "type": "stdin"
32   },
33   "input": { #控制台标准输入
34     "type": "stdin"
35   },
36   "beat": { #beat版本以及主机信息
37     "name": "itcast01",
38     "hostname": "itcast01",
39     "version": "6.5.4"
40   },
41   "host": {
42     "name": "itcast01"
43   }
44 }
45
```

2.3、读取文件

```
1 #配置读取文件项 haoke-log.yml
2
3 filebeat.inputs:
4 - type: log
5   enabled: true
6   paths:
7     - /haoke/beats/logs/*.log
8 setup.template.settings:
9   index.number_of_shards: 3
10 output.console:
11   pretty: true
12   enable: true
13
14 #启动filebeat
15 ./filebeat -e -c haoke-log.yml
16
17 #/haoke/beats/logs下创建a.log文件，并输入如下内容
18 hello
19 world
20
21 #观察filebeat输出
22 {
23   "@timestamp": "2019-01-12T14:16:10.192z",
```



```
24  "@metadata": {
25      "beat": "filebeat",
26      "type": "doc",
27      "version": "6.5.4"
28  },
29  "host": {
30      "name": "itcast01"
31  },
32  "source": "/haoke/beats/logs/a.log",
33  "offset": 0,
34  "message": "hello",
35  "prospector": {
36      "type": "log"
37  },
38  "input": {
39      "type": "log"
40  },
41  "beat": {
42      "version": "6.5.4",
43      "name": "itcast01",
44      "hostname": "itcast01"
45  }
46  }
47  {
48      "@timestamp": "2019-01-12T14:16:10.192Z",
49      "@metadata": {
50          "beat": "filebeat",
51          "type": "doc",
52          "version": "6.5.4"
53      },
54      "prospector": {
55          "type": "log"
56      },
57      "input": {
58          "type": "log"
59      },
60      "beat": {
61          "version": "6.5.4",
62          "name": "itcast01",
63          "hostname": "itcast01"
64      },
65      "host": {
66          "name": "itcast01"
67      },
68      "source": "/haoke/beats/logs/a.log",
69      "offset": 6,
70      "message": "world"
71  }
72  }
```

可以看出，已经检测到日志文件有更新，立刻就会读取到更新的内容，并且输出到控制台。

2.4、自定义字段

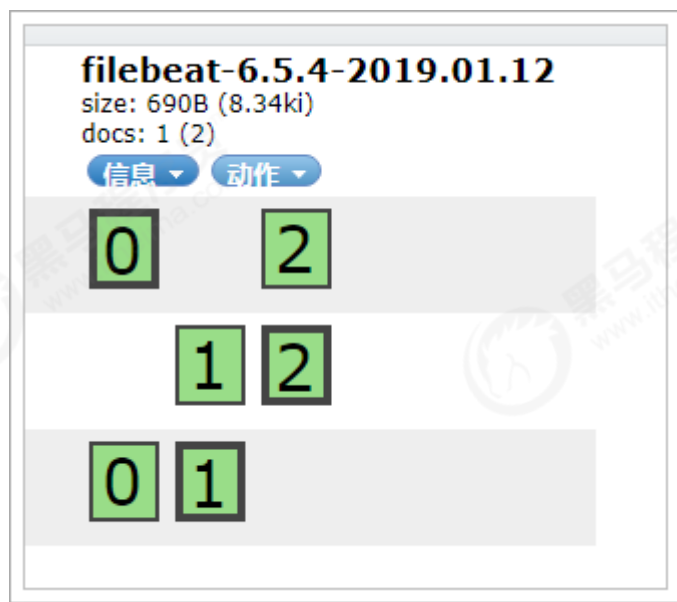


```
1 #配置读取文件项 haoke-log.yml
2
3 filebeat.inputs:
4 - type: log
5   enabled: true
6   paths:
7     - /haoke/beats/logs/*.log
8   tags: ["web"] #添加自定义tag, 便于后续的处理
9   fields: #添加自定义字段
10     from: haoke-im
11     fields_under_root: true #true为添加到根节点, false为添加到子节点中
12 setup.template.settings:
13   index.number_of_shards: 3
14 output.console:
15   pretty: true
16   enable: true
17
18 #启动filebeat
19 ./filebeat -e -c haoke-log.yml
20
21 #/haoke/beats/logs下创建a.log文件, 并输入如下内容
22 123
23
24 #执行效果
25 {
26   "@timestamp": "2019-01-12T14:37:19.845z",
27   "@metadata": {
28     "beat": "filebeat",
29     "type": "doc",
30     "version": "6.5.4"
31   },
32   "offset": 0,
33   "tags": [
34     "haoke-im"
35   ],
36   "prospector": {
37     "type": "log"
38   },
39   "beat": {
40     "name": "itcast01",
41     "hostname": "itcast01",
42     "version": "6.5.4"
43   },
44   "host": {
45     "name": "itcast01"
46   },
47   "source": "/haoke/beats/logs/a.log",
48   "message": "123",
49   "input": {
50     "type": "log"
51   },
52   "from": "haoke-im"
53 }
```

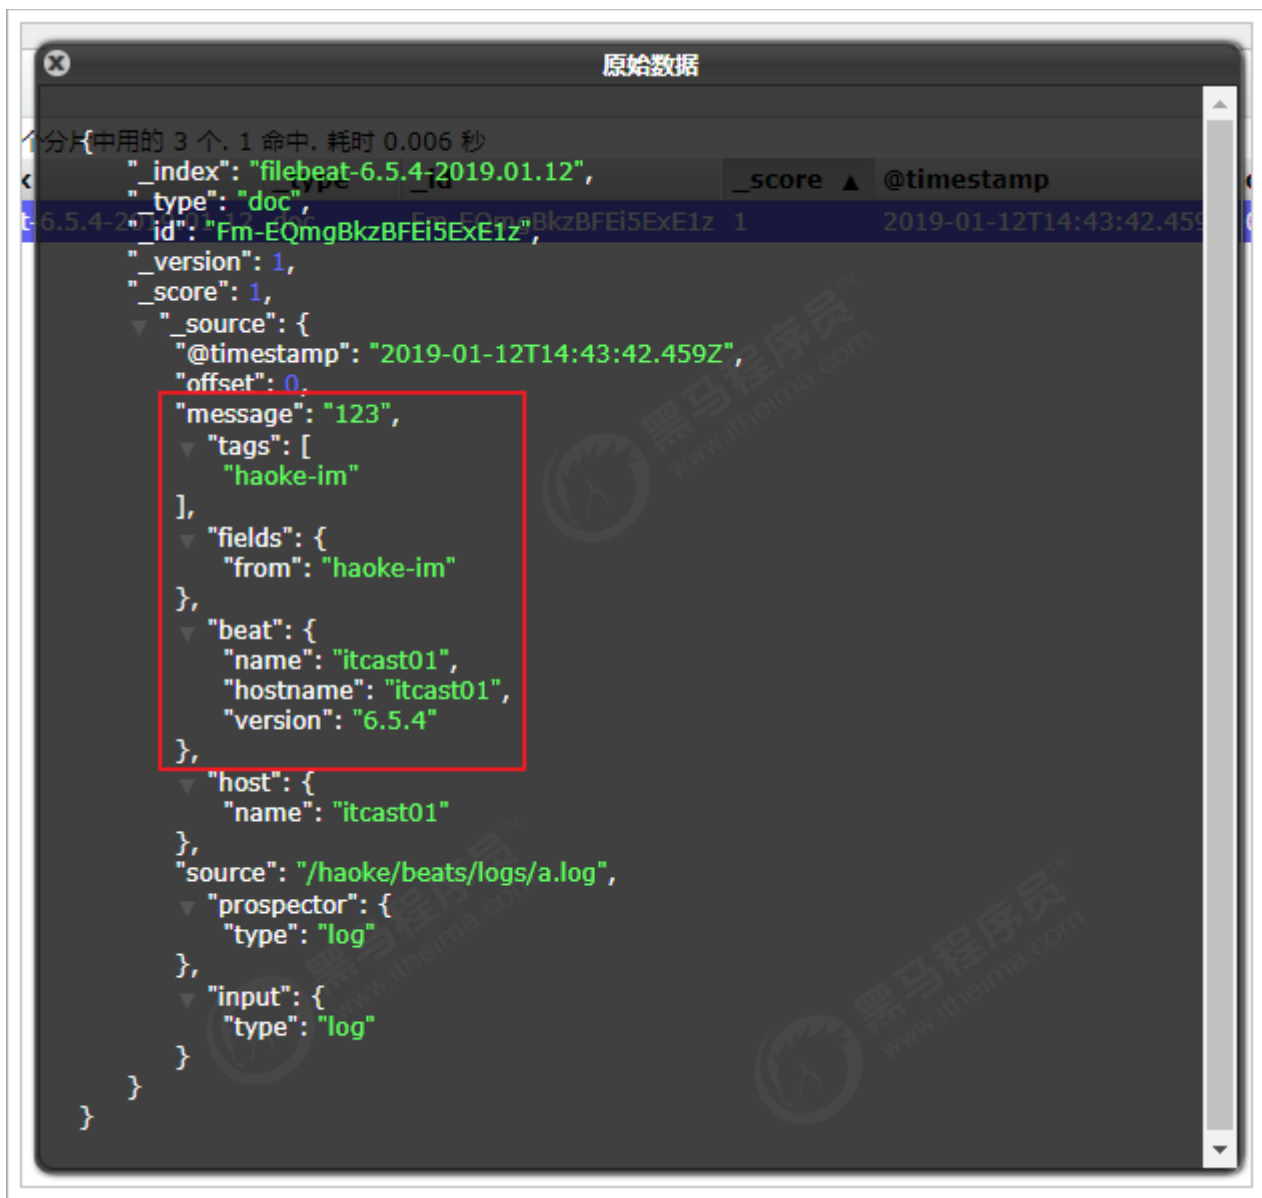
2.5、输出到Elasticsearch

```
1 # haoke-log.yml
2 filebeat.inputs:
3   - type: log
4     enabled: true
5     paths:
6       - /haoke/beats/logs/*.log
7     tags: ["haoke-im"]
8     fields:
9       from: haoke-im
10    fields_under_root: false
11  setup.template.settings:
12    index.number_of_shards: 3 #指定索引的分区数
13  output.elasticsearch: #指定ES的配置
14    hosts: ["192.168.1.7:9200", "192.168.1.7:9201", "192.168.1.7:9202"]
```

在日志文件中输入新的内容进行测试：



查看数据：



2.6、Filebeat工作原理

Filebeat由两个主要组件组成：prospector 和 harvester。

- harvester :
 - 负责读取单个文件的内容。
 - 如果文件在读取时被删除或重命名，Filebeat将继续读取文件。
- prospector
 - prospector 负责管理harvester并找到所有要读取的文件来源。
 - 如果输入类型为日志，则查找器将查找路径匹配的所有文件，并为每个文件启动一个harvester。
 - Filebeat目前支持两种prospector类型：log和stdin。
- Filebeat如何保持文件的状态
 - Filebeat 保存每个文件的状态并经常将状态刷新到磁盘上的注册文件中。
 - 该状态用于记住harvester正在读取的最后偏移量，并确保发送所有日志行。
 - 如果输出（例如Elasticsearch或Logstash）无法访问，Filebeat会跟踪最后发送的行，并在输出再次可用时继续读取文件。

- 在Filebeat运行时，每个prospector内存中也会保存的文件状态信息，当重新启动Filebeat时，将使用注册文件的数据来重建文件状态，Filebeat将每个harvester在从保存的最后偏移量继续读取。
- 文件状态记录在data/registry文件中。

启动命令：

```
1 ./filebeat -e -c haoke.yml
2 ./filebeat -e -c haoke.yml -d "publish"
3
4 #参数说明
5 -e: 输出到标准输出，默认输出到syslog和logs下
6 -c: 指定配置文件
7 -d: 输出debug信息
8
9 #测试: ./filebeat -e -c haoke-log.yml -d "publish"
10 DEBUG [publish] pipeline/processor.go:308 Publish event: {
11   "@timestamp": "2019-01-12T15:03:50.820Z",
12   "@metadata": {
13     "beat": "filebeat",
14     "type": "doc",
15     "version": "6.5.4"
16   },
17   "offset": 0,
18   "tags": [
19     "haoke-im"
20   ],
21   "input": {
22     "type": "log"
23   },
24   "prospector": {
25     "type": "log"
26   },
27   "beat": {
28     "name": "itcast01",
29     "hostname": "itcast01",
30     "version": "6.5.4"
31   },
32   "source": "/haoke/beats/logs/a.log",
33   "fields": {
34     "from": "haoke-im"
35   },
36   "host": {
37     "name": "itcast01"
38   },
39   "message": "456"
40 }
41
```

2.7、Module

前面要想实现日志数据的读取以及处理都是自己手动配置的，其实，在Filebeat中，有大量的Module，可以简化我们的配置，直接就可以使用，如下：

```
1 ./filebeat modules list
2
3 Enabled:
4
5 Disabled:
6 apache2
7 auditd
8 elasticsearch
9 haproxy
10 icinga
11 iis
12 kafka
13 kibana
14 logstash
15 mongodb
16 mysql
17 nginx
18 osquery
19 postgresql
20 redis
21 suricata
22 system
23 traefik
24
```

可以看到，内置了很多的module，但是都没有启用，如果需要启用需要进行enable操作：

```
1 ./filebeat modules enable redis #启动
2 ./filebeat modules disable redis #禁用
3
4 Enabled:
5 redis
6
7 Disabled:
8 apache2
9 auditd
10 elasticsearch
11 haproxy
12 icinga
13 iis
14 kafka
15 kibana
16 logstash
17 mongodb
18 mysql
19 nginx
20 osquery
21 postgresql
22 suricata
23 system
24 traefik
```

可以发现，redis的module已经被启用。

2.7.1、redis module

module目录：

```
1  .
2  ├── log #日志
3  │   ├── config
4  │   │   └── log.yml
5  │   ├── ingest
6  │   │   └── pipeline.json
7  │   └── manifest.yml
8  ├── module.yml
9  └── slowlog #慢查询日志
    ├── config
    │   └── slowlog.yml
    ├── ingest
    │   └── pipeline.json
    └── manifest.yml
```

2.7.2、redis module 配置

```
1  cd modules.d/
2  vim redis.yml
3
4  - module: redis
5    # Main logs
6    log:
7      enabled: true
8
9    # Set custom paths for the log files. If left empty,
10    # Filebeat will choose the paths depending on your OS.
11    var.paths: ["/data/redis-data/node01/*.log"]
12
13    # Slow logs, retrieved via the Redis API (SLOWLOG)
14    slowlog:
15      enabled: false
16
17    # The Redis hosts to connect to.
18    #var.hosts: ["localhost:6379"]
19
20    # Optional, the password to use when connecting to Redis.
21    #var.password:
```

2.7.3、修改redis的docker容器

redis默认情况下，是不会输出日志的，需要进行配置，前面我们使用的容器都没有配置日志输出，下面需要配置一下。



```
1 docker create --name redis-node01 -v /data/redis-data/node01:/data -p 6379:6379
  redis:5.0.2 --cluster-enabled yes --cluster-config-file nodes-node-01.conf --loglevel
  debug --logfile nodes-node-01.log
2
3 docker create --name redis-node02 -v /data/redis-data/node02:/data -p 6380:6379
  redis:5.0.2 --cluster-enabled yes --cluster-config-file nodes-node-02.conf --loglevel
  debug --logfile nodes-node-02.log
4
5 docker create --name redis-node03 -v /data/redis-data/node03:/data -p 6381:6379
  redis:5.0.2 --cluster-enabled yes --cluster-config-file nodes-node-03.conf --loglevel
  debug --logfile nodes-node-03.log
```

loglevel 日志等级分为：debug、verbose、notice、warning

其中，debug 会有大量信息，对开发、测试有用；

verbose 等于log4j 中的info，有很多信息，但是不会像debug那样乱；

notice 一般信息；

warning 只有非常重要/关键的消息被记录。

2.7.4、配置filebeat

```
1 #vim haoke-redis.yml
2
3 filebeat.inputs:
4 - type: log
5   enabled: true
6   paths:
7     - /haoke/log/*.log
8 setup.template.settings:
9   index.number_of_shards: 3
10 output.console:
11   pretty: true
12   enable: true
13 filebeat.config.modules:
14   path: ${path.config}/modules.d/*.yaml
15   reload.enabled: false
16
```

2.7.5、测试

```
1 ./filebeat -e -c haoke-redis.yml --modules redis
```

查询 3 个分片中用的 3 个, 60397 命中, 耗时 0.010 秒

_index	_type	_id	_score ▲	offset	prospector.type	source	file
filebeat-6.5.4-2019.01.13	doc	pm_NQmgBkzBFEi5Eil-m	1	1390083	log	/data/redis-data/node01/nodes-node-01.log	r ^
filebeat-6.5.4-2019.01.13	doc	q2_NQmgBkzBFEi5Eil-m	1	1390589	log	/data/redis-data/node01/nodes-node-01.log	r
filebeat-6.5.4-2019.01.13	doc	r2_NQmgBkzBFEi5Eil-m	1	1390981	log	/data/redis-data/node01/nodes-node-01.log	r
filebeat-6.5.4-2019.01.13	doc	s2_NQmgBkzBFEi5Eil-m	1	1391373	log	/data/redis-data/node01/nodes-node-01.log	r
filebeat-6.5.4-2019.01.13	doc	tG_NQmgBkzBFEi5Eil-m	1	1391455	log	/data/redis-data/node01/nodes-node-01.log	r
filebeat-6.5.4-2019.01.13	doc	u2_NQmgBkzBFEi5Eil-m	1	1392157	log	/data/redis-data/node01/nodes-node-01.log	r
filebeat-6.5.4-2019.01.13	doc	vG_NQmgBkzBFEi5Eil-m	1	1392239	log	/data/redis-data/node01/nodes-node-01.log	r
filebeat-6.5.4-2019.01.13	doc	vW_NQmgBkzBFEi5Eil-m	1	1392321	log	/data/redis-data/node01/nodes-node-01.log	r
filebeat-6.5.4-2019.01.13	doc	wW_NQmgBkzBFEi5Eil-m	1	1392713	log	/data/redis-data/node01/nodes-node-01.log	r
filebeat-6.5.4-2019.01.13	doc	wm_NQmgBkzBFEi5Eil-m	1	1392827	log	/data/redis-data/node01/nodes-node-01.log	r
filebeat-6.5.4-2019.01.13	doc	xW_NQmgBkzBFEi5Eil-m	1	1393105	log	/data/redis-data/node01/nodes-node-01.log	r
filebeat-6.5.4-2019.01.13	doc	x2_NQmgBkzBFEi5Eil-m	1	1393333	log	/data/redis-data/node01/nodes-node-01.log	r
filebeat-6.5.4-2019.01.13	doc	yW_NQmgBkzBFEi5Eil-m	1	1393497	log	/data/redis-data/node01/nodes-node-01.log	r

测试发现，数据已经写入到了Elasticsearch中。

当然了，其他的Module的用法参加官方文档：

<https://www.elastic.co/guide/en/beats/filebeat/current/filebeat-modules.html>

- [Modules overview](#)
- [Apache2 module](#)
- [Auditd module](#)
- [Elasticsearch module](#)
- [haproxy module](#)
- [Icinga module](#)
- [IIS module](#)
- [Kafka module](#)
- [Kibana module](#)
- [Logstash module](#)
- [MongoDB module](#)
- [MySQL module](#)
- [Nginx module](#)
- [Osquery module](#)
- [PostgreSQL module](#)
- [Redis module](#)
- [Suricata module](#)
- [System module](#)

3、Metricbeat



Metricbeat

轻量型指标采集器

用于从系统和服務收集指标。Metricbeat 能够以一种轻量型的方式，输送各种系统和服務统计数据，从 CPU 到内存，从 Redis 到 Nginx，不一而足。

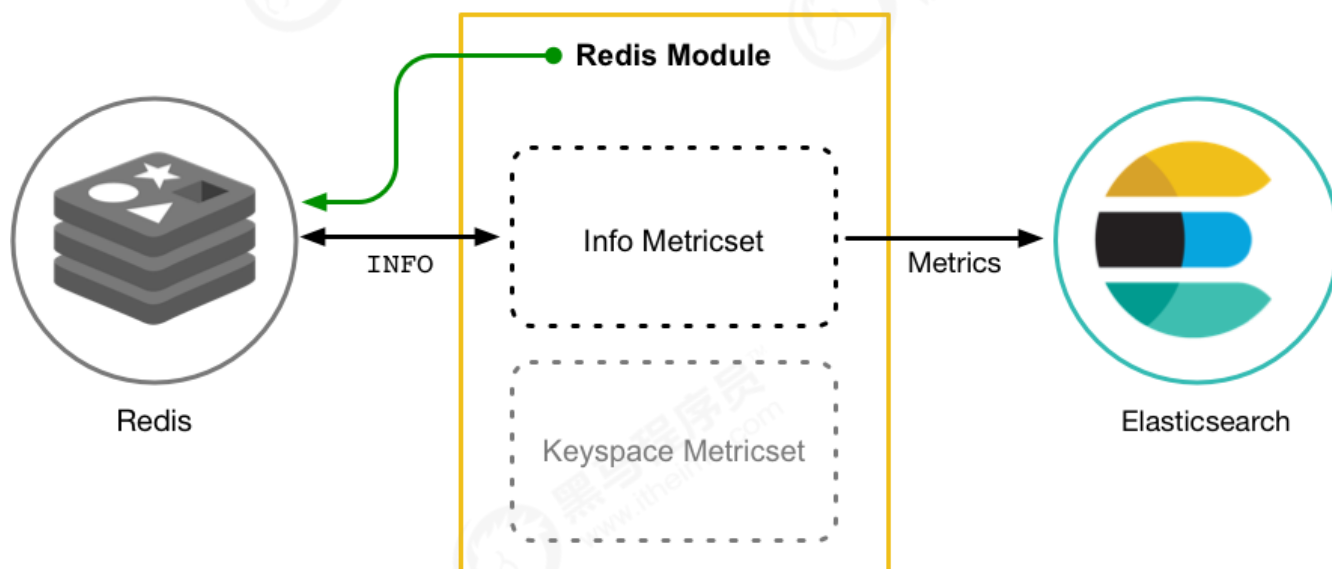
- 定期收集操作系统或应用服务的指标数据
- 存储到Elasticsearch中，进行实时分析

3.1、Metricbeat组成

Metricbeat有2部分组成，一部分是Module，另一部分为Metricset。

- Module
 - 收集的对象，如：mysql、redis、操作系统等；
- Metricset
 - 收集指标的集合，如：cpu、memory、network等；

以Redis Module为例：



3.2、部署与收集系统指标

```
1 tar -xvf metricbeat-6.5.4-linux-x86_64.tar.gz
2 cd metricbeat-6.5.4-linux-x86_64
3 vim metricbeat.yml
4
5 metricbeat.config.modules:
```



```

6   path: ${path.config}/modules.d/*.yaml
7   reload.enabled: false
8   setup.template.settings:
9     index.number_of_shards: 1
10    index.codec: best_compression
11  setup.kibana:
12  output.elasticsearch:
13    hosts: ["192.168.1.7:9200", "192.168.1.7:9201", "192.168.1.7:9202"]
14  processors:
15    - add_host_metadata: ~
16    - add_cloud_metadata: ~
17
18  #启动
19  ./metricbeat -e
  
```

在Elasticsearch中可以看到，系统的一些指标数据已经写入进去了：

查询 1 个分片中用的 1 个, 94 命中, 耗时 0.009 秒

_index	_type	_id	_score	@timestamp	metricset.name	metricset.module	me
metricbeat-6.5.4-2019.01.13	doc	etSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.306Z	uptime	system	
metricbeat-6.5.4-2019.01.13	doc	e9SLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.306Z	fsstat	system	
metricbeat-6.5.4-2019.01.13	doc	fNSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.306Z	filesystem	system	
metricbeat-6.5.4-2019.01.13	doc	fdSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.306Z	filesystem	system	
metricbeat-6.5.4-2019.01.13	doc	ftSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.306Z	filesystem	system	
metricbeat-6.5.4-2019.01.13	doc	f9SLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.306Z	filesystem	system	
metricbeat-6.5.4-2019.01.13	doc	gNSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.306Z	filesystem	system	
metricbeat-6.5.4-2019.01.13	doc	gdSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.306Z	filesystem	system	
metricbeat-6.5.4-2019.01.13	doc	gtSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.306Z	filesystem	system	
metricbeat-6.5.4-2019.01.13	doc	g9SLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.306Z	filesystem	system	
metricbeat-6.5.4-2019.01.13	doc	hNSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.306Z	filesystem	system	
metricbeat-6.5.4-2019.01.13	doc	hdSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.307Z	cpu	system	
metricbeat-6.5.4-2019.01.13	doc	htSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.307Z	load	system	
metricbeat-6.5.4-2019.01.13	doc	h9SLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.307Z	memory	system	
metricbeat-6.5.4-2019.01.13	doc	iNSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.307Z	network	system	
metricbeat-6.5.4-2019.01.13	doc	idSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.307Z	network	system	
metricbeat-6.5.4-2019.01.13	doc	itSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.307Z	network	system	
metricbeat-6.5.4-2019.01.13	doc	i9SLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.307Z	network	system	
metricbeat-6.5.4-2019.01.13	doc	jNSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.307Z	network	system	
metricbeat-6.5.4-2019.01.13	doc	jdSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.307Z	network	system	
metricbeat-6.5.4-2019.01.13	doc	jtSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.307Z	network	system	
metricbeat-6.5.4-2019.01.13	doc	j9SLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.307Z	network	system	
metricbeat-6.5.4-2019.01.13	doc	kNSLRmgB8SxFt2LBLp1y	1	2019-01-13T09:29:11.307Z	network	system	

system module配置：

```

1  root@itcast01:modules.d# cat system.yml
2  # Module: system
3  # Docs: https://www.elastic.co/guide/en/beats/metricbeat/6.5/metricbeat-module-
    system.html
4
5  - module: system
6    period: 10s
7    metricsets:
8      - cpu
9      - load
10     - memory
11     - network
12     - process
  
```



```
13 - process_summary
14   #- core
15   #- diskio
16   #- socket
17   process.include_top_n:
18     by_cpu: 5      # include top 5 processes by CPU
19     by_memory: 5   # include top 5 processes by memory
20
21 - module: system
22   period: 1m
23   metricsets:
24     - filesystem
25     - fsstat
26   processors:
27     - drop_event.when.regex:
28       system.filesystem.mount_point: '^/(sys|cgroup|proc|dev|etc|host|lib)($|/)'
29
30 - module: system
31   period: 15m
32   metricsets:
33     - uptime
34
35   #- module: system
36   # period: 5m
37   # metricsets:
38   #   - raid
39   # raid.mount_point: '/'
```

3.3、Module

```
1 ./metricbeat modules list #查看列表
2
3 Enabled:
4 system #默认启用
5
6 Disabled:
7 aerospike
8 apache
9 ceph
10 couchbase
11 docker
12 dropwizard
13 elasticsearch
14 envoyproxy
15 etcd
16 golang
17 graphite
18 haproxy
19 http
20 jolokia
21 kafka
22 kibana
```

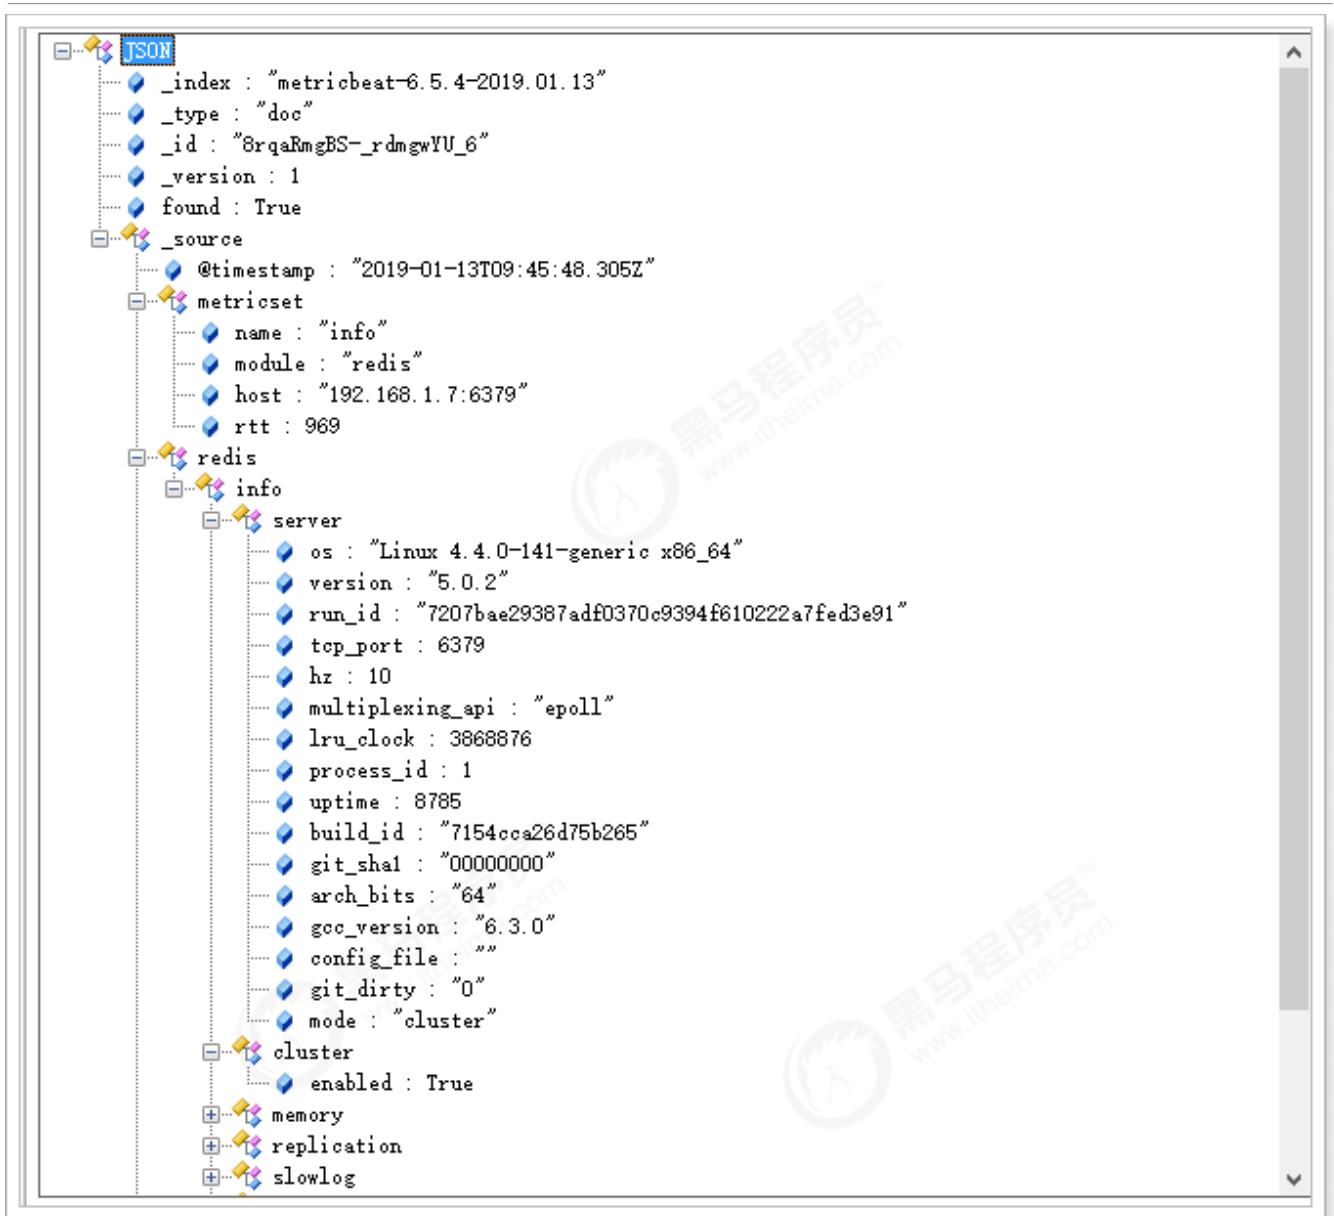
```
23 kubernetes
24 kvm
25 logstash
26 memcached
27 mongodb
28 munin
29 mysql
30 nginx
31 php_fpm
32 postgresql
33 prometheus
34 rabbitmq
35 redis
36 traefik
37 uwsgi
38 vsphere
39 windows
40 zookeeper
```

3.4、Redis Module

```
1 #启用redis module
2 ./metricbeat modules enable redis
3 #修改redis module配置
4 vim modules.d/redis.yml
5
6 - module: redis
7   metricsets:
8     - info
9     # - keyspace
10    period: 10s
11
12    # Redis hosts
13    hosts: ["192.168.1.7:6379","192.168.1.7:6380","192.168.1.7:6381"]
14
15    # Network type to be used for redis connection. Default: tcp
16    #network: tcp
17
18    # Max number of concurrent connections. Default: 10
19    #maxconn: 10
20
21    # Redis AUTH password. Empty by default.
22    #password: foobared
23
24 #启动
25 ./metricbeat -e
```

可以看到Redis的指标数据写入到了Elasticsearch中：

metricbeat-6.5.4-2019.01.13	doc	8rqARmgBS-_rdmgwYU_6	1	2019-01-13T09:45:48.305Z	969	info	redis
metricbeat-6.5.4-2019.01.13	doc	87qaRmgBS-_rdmgwYU_6	1	2019-01-13T09:45:48.305Z	2269	info	redis
metricbeat-6.5.4-2019.01.13	doc	9LqaRmgBS-_rdmgwYU_6	1	2019-01-13T09:45:48.310Z	581	info	redis



更多的Module使用参见官方文档：

<https://www.elastic.co/guide/en/beats/metricbeat/current/metricbeat-modules.html>

4、Kibana



您使用 Elastic Stack 的窗口

通过 Kibana，您能够对 Elasticsearch 中的数据进行可视化并在 Elastic Stack 进行操作，因此您可以在这里解开任何疑问：例如，为何会在凌晨 2:00 收到传呼，雨水会对季度数据造成怎样的影响。

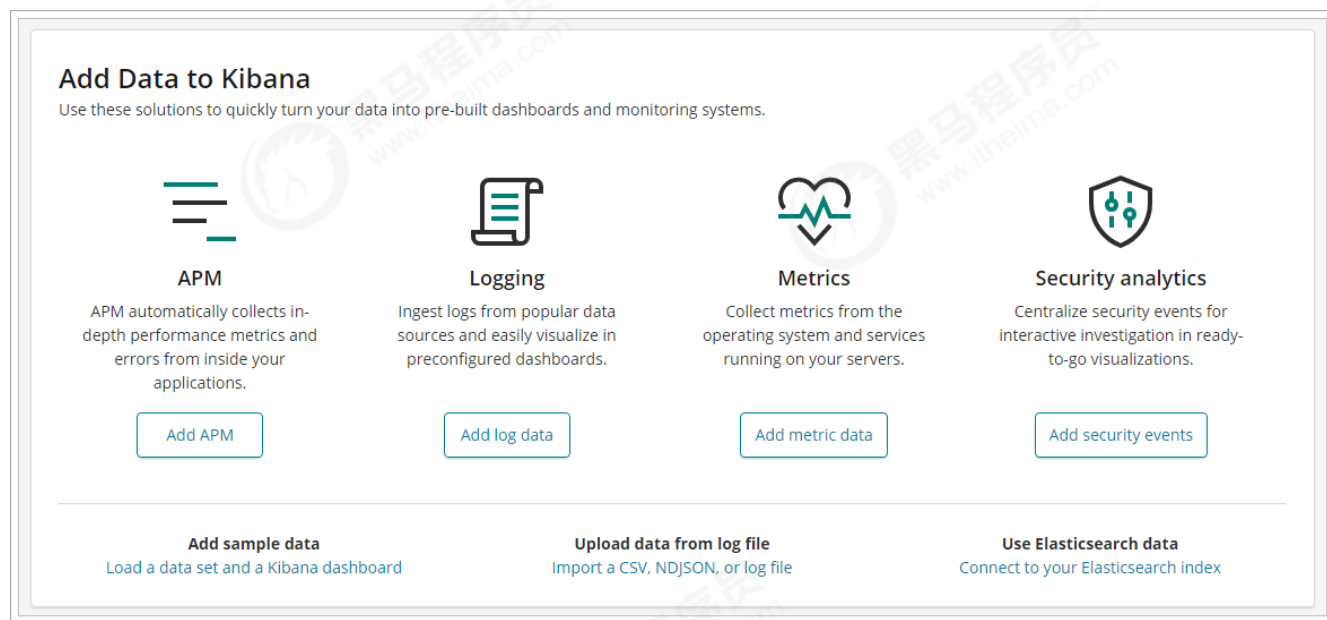


Kibana 是一款开源的数据分析和可视化平台，它是 Elastic Stack 成员之一，设计用于和 Elasticsearch 协作。您可以使用 Kibana 对 Elasticsearch 索引中的数据进行搜索、查看、交互操作。您可以很方便的利用图表、表格及地图对数据进行多元化的分析和呈现。

官网：<https://www.elastic.co/cn/products/kibana>

4.1、配置安装

```
1 #解压安装包
2 tar -xvf kibana-6.5.4-linux-x86_64.tar.gz
3
4 #修改配置文件
5 vim config/kibana.yml
6
7 server.host: "192.168.1.7" #对外暴露服务的地址
8 elasticsearch.url: "http://192.168.1.7:9200" #配置Elasticsearch
9
10 #启动
11 ./bin/kibana
12
13 #通过浏览器进行访问
14 http://192.168.1.7:5601/app/kibana
```



可以看到kibana页面，并且可以看到提示，导入数据到Kibana。

4.2、通过docker部署

```
1 #拉取镜像
2 docker pull kibana:6.5.4
3
4 #创建配置文件
5 vim kibana.yml
6 server.host: "192.168.1.7"
7 elasticsearch.url: "http://192.168.1.7:9200"
8
9 #创建容器
10 docker create --name kibana --net host -v /haoke/beats/kibana-
  docker/kibana.yml:/usr/share/kibana/config/kibana.yml kibana:6.5.4
11
12 #启动容器
13 docker logs -f kibana
```

4.3、功能说明



4.4、数据探索

首先先添加索引信息：

The screenshot shows the Kibana Management interface. On the left sidebar, the 'Management' tab is selected. The main content area shows the 'Index Patterns' section with a 'Create index pattern' button. Below this, the index pattern 'metricbeat-*' is selected. On the right, the 'Time Filter field name: @timestamp' is highlighted. Below this, a table lists fields and their types:

Name	Type	Format
@timestamp	date	
_id	string	
_index	string	
_score	number	

即可查看索引数据：

The screenshot shows the Kibana Discover interface. At the top, it displays '24 hits' and a search bar with the query 'status:200 AND extension:PHP'. Below the search bar, there is a histogram showing the distribution of data over time. The x-axis is labeled '@timestamp per 20 milliseconds' and the y-axis is labeled 'Count'. Below the histogram, there is a table of search results with columns for 'Time' and '_source'.

Time	_source
January 14th 2019, 15:58:18.393	@timestamp: January 14th 2019, 15:58:18.393 host.name: itcast01 host.containerized: false host.architecture: x86_64 host.os.family: debian host.os.codename: xenial host.os.platform: ubuntu host.os.version: 16.04.5 LTS (Xenial Xerus) host.id: 92508adbe358027654289985c356d36 metricset.name: process metricset.module: system metricset.rtt: 64,399 system.process.cmdline: /opt/jdk-11.0.1/bin/java -Xms512m -Xmx512m -XX:+UseConcMarkSweepGC -XX:CMSInitiatingOccupancyFraction=75 -X
January 14th 2019, 15:58:18.393	@timestamp: January 14th 2019, 15:58:18.393 metricset.name: process metricset.module: system metricset.rtt: 64,416 system.process.state: sleeping system.process.cpu.total.value: 267,840 system.process.cpu.total.pct: 0.01 system.process.cpu.total.norm.pct: 0.003 system.process.cpu.start_time: January 14th 2019, 09:25:27.000 system.process.ppid: 2,918

4.5、Metricbeat 仪表盘

可以将Metricbeat的数据在Kibana中展示。



```
1 #修改metricbeat配置
2 setup.kibana:
3   host: "192.168.1.7:5601"
4
5 #安装仪表盘到kibana
6 ./metricbeat setup --dashboards
7
```

即可在Kibana中看到仪表盘数据：

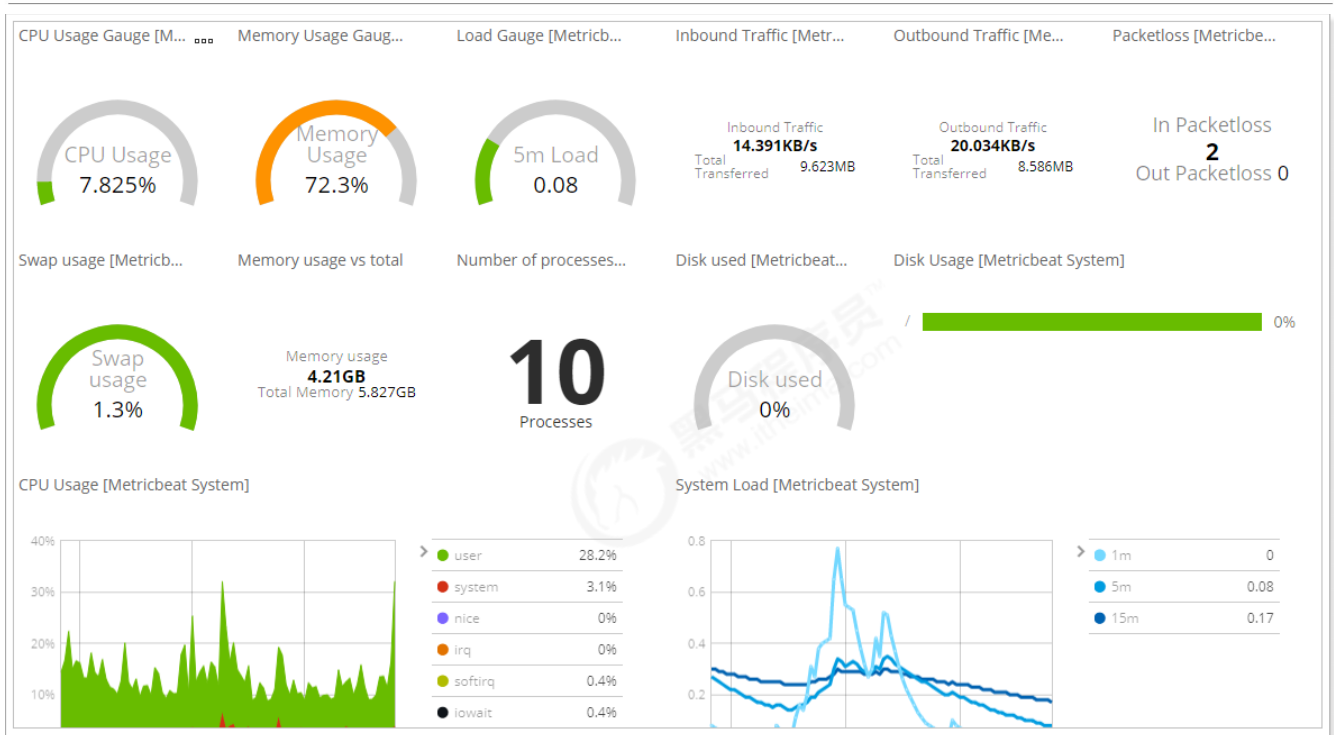
Dashboards

[Create new dashboard](#)

Q Search...

<input type="checkbox"/>	Title	Description	Actions
<input type="checkbox"/>	[Metricbeat Docker] Overview	Overview of docker containers	Edit
<input type="checkbox"/>	[Metricbeat Apache] Overview	Overview of Apache server status	Edit
<input type="checkbox"/>	[Metricbeat System] Containers overview	Overview of container metrics	Edit
<input type="checkbox"/>	[Metricbeat System] Host overview	Overview of host metrics	Edit
<input type="checkbox"/>	[Metricbeat Kafka] Overview	Kafka analysis of topics and consumer groups	Edit
<input type="checkbox"/>	[Metricbeat Golang] Overview	Overview of Go profiling information	Edit
<input type="checkbox"/>	[Metricbeat HAProxy] HTTP backend	HAProxy HTTP backend metrics	Edit
<input type="checkbox"/>	[Metricbeat HAProxy] HTTP server	HAProxy metrics for HTTP mode	Edit
<input type="checkbox"/>	[Metricbeat HAProxy] Backend	HAProxy backend metrics	Edit
<input type="checkbox"/>	[Metricbeat HAProxy] Overview	HAProxy overview	Edit
<input type="checkbox"/>	[Metricbeat HAProxy] HTTP frontend	HAProxy frontend metrics	Edit

查看系统信息：



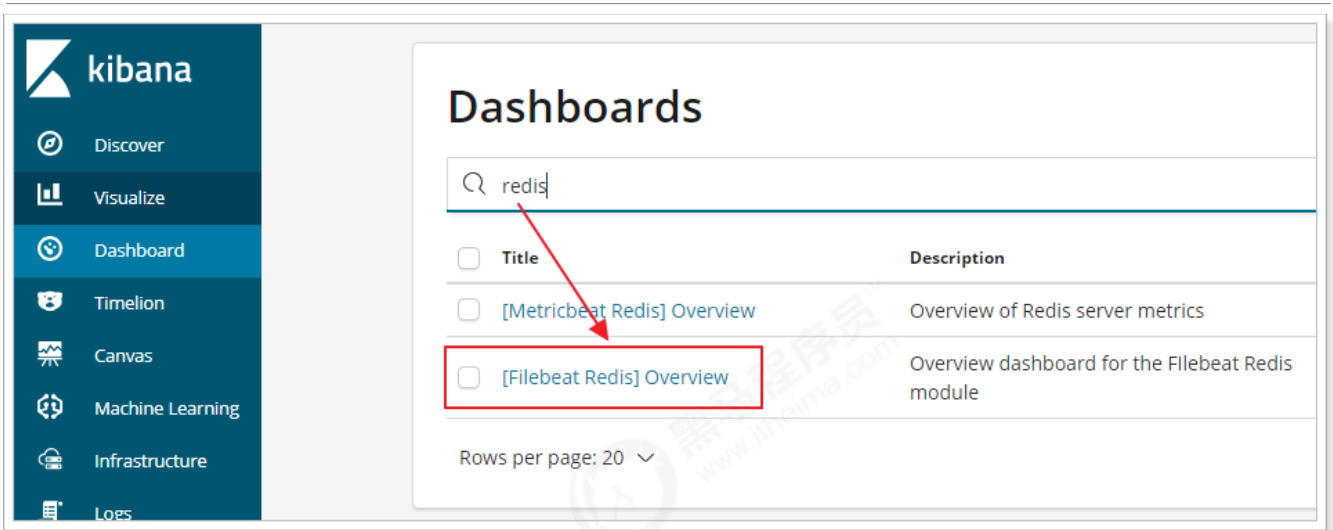
4.6、Filebeat 仪表盘

以Redis为例：

```

1  #修改配置文件
2  filebeat.inputs:
3  - type: log
4    enabled: true
5    paths:
6      - /haoke/log/*.log
7  setup.template.settings:
8    index.number_of_shards: 3
9  filebeat.config.modules:
10   path: ${path.config}/modules.d/*.yaml
11   reload.enabled: false
12  output.elasticsearch:
13   hosts: ["192.168.1.7:9200", "192.168.1.7:9201", "192.168.1.7:9202"]
14  setup.kibana:
15   host: "192.168.1.7:5601"
16
17  #安装仪表盘到kibana
18  ./filebeat -c haoke-redis.yaml setup

```

在kibana中已经看到了Filebeat中的redis仪表盘。

1-50 of 767,123				
Time	beat.name	redis.log.level	redis.log.role	redis.log.message
▶ January 14th 2019, 17:21:54.047	itcast01	-	-	-
▶ January 14th 2019, 17:21:53.046	itcast01	-	-	-
▶ January 14th 2019, 17:21:53.046	itcast01	-	-	-
▶ January 14th 2019, 17:21:53.046	itcast01	-	-	-
▶ January 14th 2019, 17:21:53.046	itcast01	-	-	-
▶ January 14th 2019, 17:21:53.046	itcast01	-	-	-
▶ January 14th 2019, 17:21:53.046	itcast01	-	-	-

5、Logstash

5.1、简介




集中、转换和存储数据

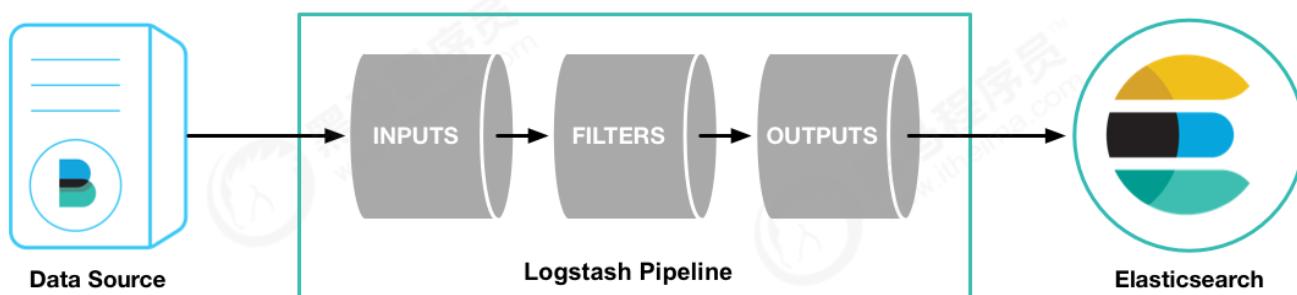
Logstash 是开源的服务器端数据处理管道，能够同时从多个来源采集数据，转换数据，然后将数据发送到您最喜欢的“存储库”中。（我们的存储库当然是 Elasticsearch。）



用途：



5.2、部署安装




```
1 #检查jdk环境, 要求jdk1.8+
2 java -version
3
4 #解压安装包
5 tar -xvf logstash-6.5.4.tar.gz
6
7 #第一个logstash示例
8 bin/logstash -e 'input { stdin { } } output { stdout { } }'
```

执行效果如下：

```
[2019-01-14T18:17:44,254][INFO ][logstash.agent
nt {:port=>9600}
hello
{
  "message" => "hello",
  "@timestamp" => 2019-01-14T10:17:51.638Z,
  "host" => "itcast01",
  "@version" => "1"
}
```

5.3、接收Filebeat输入的日志

接下来，我们将Filebeat和Logstash整合起来，读取nginx的日志。



5.3.1、安装Nginx

```
1 apt install nginx -y
2
3 #/usr/sbin/nginx : 主程序
4 #/etc/nginx : 存放配置文件
5 #/usr/share/nginx : 存放静态文件
6 #/var/log/nginx : 存放日志
7
8 #nginx服务命令
9 service nginx {start|stop|restart|reload|force-
  reload|status|configtest|rotate|upgrade}
10
11 #通过浏览器访问页面并且查看日志
12 #访问地址 : http://192.168.1.7/
13 tail -f /var/log/nginx/access.log
```



```
root@itcast01:haoke# tail -f /var/log/nginx/access.log
192.168.1.20 - - [14/Jan/2019:19:52:49 +0800] "GET / HTTP/1.1" 200 396 "-" "Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/70.0.3538.67 Safari/537.36"
192.168.1.20 - - [14/Jan/2019:19:52:49 +0800] "GET /favicon.ico HTTP/1.1" 404 209 "http://192.168.1.7/" "Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/70.0.3538.67 Safari/537.36"

192.168.1.20 - - [14/Jan/2019:19:53:21 +0800] "GET / HTTP/1.1" 304 0 "-" "Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/70.0.3538.67 Safari/537.36"
```

5.3.2、配置Filebeat

```
1 #vim haoke-nginx.yml
2 filebeat.inputs:
3 - type: log
4   enabled: true
5   paths:
6     - /var/log/nginx/access.log
7   tags: ["log"]
8   fields:
9     from: nginx
10    fields_under_root: false
11 output.logstash:
12   hosts: ["192.168.1.7:5044"]
13
14 #启动
15 ./filebeat -e -c haoke-nginx.yml
16 #说明：现在启动会报错，因为Logstash还没有启动
```

5.3.3、配置Logstash

```
1 vim haoke-pipeline.conf
2 #输入如下内容：
3 input {
4   beats {
5     port => "5044"
6   }
7 }
8
9 # The filter part of this file is commented out to indicate that it is
10 # optional.
11 # filter {
12 #
13 # }
14
15 output {
16   stdout { codec => rubydebug }
17 }
18
19 #启动 --config.test_and_exit 用于测试配置文件是否正确
20 bin/logstash -f haoke-pipeline.conf --config.test_and_exit
21
```

```
22 #[INFO ][logstash.runner          ] Using config.test_and_exit mode. Config
    Validation Result: OK. Exiting Logstash
23
24 #正式启动 --config.reload.automatic 热加载配置文件，修改配置文件后无需重新启动
25 bin/logstash -f haoke-pipeline.conf --config.reload.automatic
```

5.3.4、测试

分别启动Filebeat和Logstash，刷新页面查看输出。

```
1  {
2    "@timestamp" => 2019-01-14T12:23:37.604Z,
3    "fields" => {
4      "from" => "nginx"
5    },
6    "source" => "/var/log/nginx/access.log",
7    "tags" => [
8      [0] "log",
9      [1] "beats_input_codec_plain_applied"
10   ],
11    "host" => {
12      "name" => "itcast01"
13    },
14    "beat" => {
15      "name" => "itcast01",
16      "version" => "6.5.4",
17      "hostname" => "itcast01"
18    },
19    "@version" => "1",
20    "offset" => 600,
21    "message" => "192.168.1.20 - - [14/Jan/2019:20:23:35 +0800] \"GET / HTTP/1.1\"
    304 0 \"-\" \"Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML,
    like Gecko) Chrome/70.0.3538.67 Safari/537.36\"",
22    "input" => {
23      "type" => "log"
24    },
25    "prospector" => {
26      "type" => "log"
27    }
28  }
29
```

可以看到，已经在控制台输出了nginx的访问日志。

5.3.4、配置filter

在前面的输出中，可以看出，虽然可以拿到日志信息，但是信息格式并不友好，比如说，不能直接拿到日志中的ip地址。

第一步，自定义nginx的日志格式



```
1 vim /etc/nginx/nginx.conf
2
3 log_format main '$remote_addr - $remote_user [$time_local] '
4                 '$request' $status $body_bytes_sent '
5                 '$http_referer' '$http_user_agent';
6
7 access_log /var/log/nginx/access.log main;
8
9 nginx -s reload
```

第二步，编写nginx-patterns文件

```
1 NGINX_ACCESS %{IPORHOST:remote_addr} - %{USERNAME:remote_user} \[%
  {HTTPDATE:time_local}\] \" %{DATA:request}\" %{INT:status} %{NUMBER:bytes_sent} \"%
  {DATA:http_referer}\" \" %{DATA:http_user_agent}\"
```

第三步，修改haoke-pipeline.conf文件

```
1 input {
2     beats {
3         port => "5044"
4     }
5 }
6
7 filter {
8     grok {
9         patterns_dir => "/haoke/logstash-6.5.4/nginx-patterns"
10        match => { "message" => "%{NGINX_ACCESS}" }
11        remove_tag => [ "_grokparsefailure" ]
12        add_tag => [ "nginx_access" ]
13    }
14 }
15
16 output {
17     stdout { codec => rubydebug }
18 }
```

第四步，测试

结果：

```
1 {
2     "input" => {
3         "type" => "log"
4     },
5     "http_referer" => "-",
6     "remote_addr" => "192.168.1.20",
7     "fields" => {
8         "from" => "nginx"
9     },
10    "host" => {
```



```
11     "name" => "itcast01"
12 },
13     "offset" => 664,
14     "request" => "GET / HTTP/1.1",
15     "message" => "192.168.1.20 - - [14/Jan/2019:22:52:49 +0800] \"GET /
HTTP/1.1\" 304 0 \"-\" \"Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36
(KHTML, like Gecko) Chrome/70.0.3538.67 Safari/537.36\"",
16     "bytes_sent" => "0",
17     "status" => "304",
18     "time_local" => "14/Jan/2019:22:52:49 +0800",
19     "http_user_agent" => "Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36
(KHTML, like Gecko) Chrome/70.0.3538.67 Safari/537.36",
20     "prospector" => {
21     "type" => "log"
22 },
23     "@version" => "1",
24     "remote_user" => "-",
25     "@timestamp" => 2019-01-14T14:52:58.948Z,
26     "source" => "/var/log/nginx/access.log",
27     "beat" => {
28     "name" => "itcast01",
29     "hostname" => "itcast01",
30     "version" => "6.5.4"
31 },
32     "tags" => [
33     [0] "log",
34     [1] "beats_input_codec_plain_applied",
35     [2] "nginx_access"
36 ]
37 }
38
```

5.3.5、发送到Elasticsearch

修改配置：

```
1 #vim haoke-pipeline.conf
2 input {
3     beats {
4         port => "5044"
5     }
6 }
7
8 filter {
9     grok {
10         patterns_dir => "/haoke/logstash-6.5.4/nginx-patterns"
11         match => { "message" => "%{NGINX_ACCESS}" }
12         remove_tag => [ "_grokparsefailure" ]
13         add_tag => [ "nginx_access" ]
14     }
15 }
16
17 #output {
```



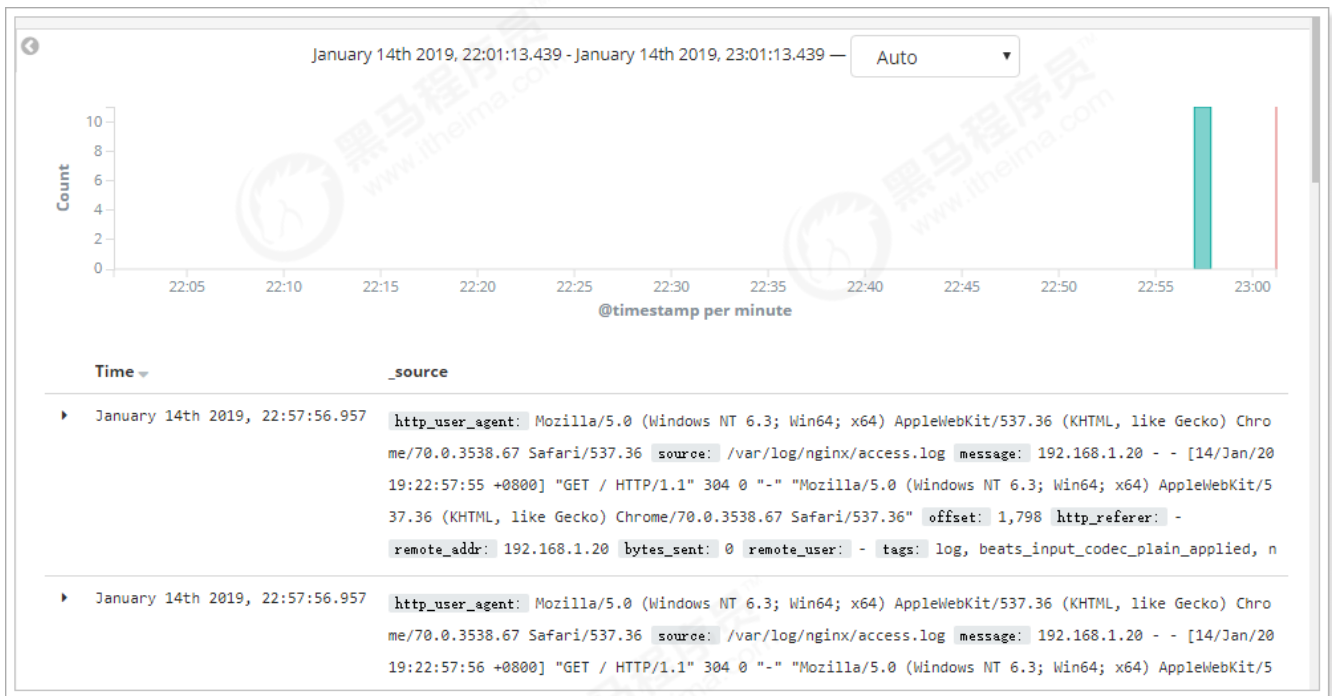
```
18 #      stdout { codec => rubydebug }
19 #}
20 output {
21   elasticsearch {
22     hosts => [ "192.168.1.7:9200", "192.168.1.7:9201", "192.168.1.7:9202" ]
23   }
24 }
```

测试：

查询 5 个分片中用的 5 个, 11 命中, 耗时 0.003 秒

_index	_type	_id	_score	http_user_agent
logstash-2019.01.14	doc	sJfeTGgBDKu_KJg5gAF5	1	Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko)
logstash-2019.01.14	doc	spfeTGgBDKu_KJg5hQFH	1	Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko)
logstash-2019.01.14	doc	JO3eTGgBDCmA2seYhTxI	1	Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko)
logstash-2019.01.14	doc	sZfeTGgBDKu_KJg5hQFH	1	Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko)
logstash-2019.01.14	doc	EzbeTGgBqy2Kgo-ShV1H	1	Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko)
logstash-2019.01.14	doc	FDbeTGgBqy2Kgo-ShV1H	1	Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko)
logstash-2019.01.14	doc	FTbeTGgBqy2Kgo-ShV1H	1	Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko)
logstash-2019.01.14	doc	Ie3eTGgBDCmA2seYhTxK	1	Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko)
logstash-2019.01.14	doc	Iu3eTGgBDCmA2seYhTxK	1	Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko)
logstash-2019.01.14	doc	I-3eTGgBDCmA2seYhTxI	1	Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko)
logstash-2019.01.14	doc	s5feTGgBDKu_KJg5hQFH	1	Mozilla/5.0 (Windows NT 6.3; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko)

在Kibana中查看：



制作柱形图：

